

Task: "turn on the stove..."


Original CoT

TASK: pick up the book...
PLAN: 1. move to the black book
2. grasp the black book
3. move the black book to the desk caddy
4. release the black book...
SUBTASK: move to the black book
MOVE: move back and left and down
GRIPPER POSITION: [49, 101]

Edited CoT (direction_flip)

TASK: pick up the book...
PLAN: 1. move to the black book
2. grasp the black book
3. move the black book to the desk caddy
4. release the black book...
SUBTASK: move to the black book
MOVE: move **forward** and **right** and **up**
GRIPPER POSITION: [49, 101]

Action delta ($\Delta_\infty=1.55$)



CoT-Faith measures 3 things per model

P1 Attention decomposition
visual / instruction / cot / action-prev

P2 Causal edit (Δ -action)
10 edit families $\times$ 100 samples

P3 Attention $\rightarrow$ error AUROC
per-sample failure predictor

action-token
attention split
(ECot-bridge)

